

Complementary notebook for ShoupModel.jl

This File

The purpose of this notebook is to accompany the ShoupModel.jl file, acting as supportive documentation for using the model. The model is based on the 2006 paper by Donal Shoup, [Cruising for Parking](#). This notebook is interactive, and I'd urge you to play around with the parameters. Also, any suggestions on further development are more than welcome!

The Paper

In the paper, Shoup(2006) presents a model which seeks to capture how the price-ratio between curb-side/off-street parking, fuel cost, and an individual's value of time impacts the incentive for cruising. The paper concludes that when curb-side parking is underpriced, it creates an incentive for individuals to cruise for parking, resulting in congestion and air pollution. Bringing the price of curb-side parking in-line with the off-street parking price can consequently yield a triple-dividend: reducing search times, reduce congestion, and raise revenue to reduce the deadweight loss from other forms taxation.

The Model

The core model, as presented in the paper, has seven variables:

- p - price of curb-side parking (\$/h)
- m - price of off-street parking (\$/h)
- t - parking duration (h)
- c - time spent searching for parking at the curb (h)
- f - fuel cost of cruising (\$/h)
- n - number of people in the car (persons)
- v - value of time spent cruising (\$/h/person)

From these variables, we can define three terms:

Equation	Definition
$t(m - p)$	money saved by parking at the curb
fc	money cost of cruising for curb parking
nvc	monetized cost of time spent cruising for curb parking
$fc + nvc = c(f + nv)$	money and (monetized) time cost of cruising for curb parking

First, the price money saved is the amount of time you wish to park multiplied by the hourly price difference between curbside and off-street parking. This quantity represents the benefit, or the potential consumer surplus if the agent immediately found a space to park on the curb. On the other hand, the agent incurs two costs when searching for parking, the fuel cost, fc , and the monetized cost of time, nvc . The combined cost of these two would be the total cost incurred by an individual who is cruising.

The equilibrium cruising time is where the cost of cruising equates the potential benefits, that's to say:

$$c^*(f + nv) = t(m - p)$$

$$c^* = \frac{t(m - p)}{f + nv}$$

From the above, one can see that cities can impose several strategies to tackle the issue of cruising for parking.

1. If $m = p$, then there will no longer be an incentive to park on the curb. This can either be achieved by increasing the curbside parking fee, or increase the amount of off-street parking such that m reduces to the same level as p .
2. Fuel taxes or emission permits could increase the cost of fuel and consequently increase the cost of cruising. In effect, this would reduce cruising however, unlikely to eliminate it, as it doesn't tackle the root cause of the issue..
3. Similarly to increases in fuel cost, policy to promote carpooling, or secondary vehicle taxes may increase n and reduce cruising times.

Julia Implementation

Based on the simple model outlined above, one can outline a basic agent based model, where agents arrive to a curbside parking location. If curbside parking is available, they will immediately park on the curb given that $p \leq m$. If no location is available, the agent will cruise for a maximum of time of c^* . If the agent has not been able to find a location to park within c^* minutes, it will park off-street. When an available parking spot on the curb opens up, all the agents which are currently cruising for parking are equally likely to occupy the available slot.

Variable list

To enable a degree of heterogeneity between agents, almost all model inputs can be provided as distributions or absolute values. Furthermore, to due to a lack of information on certain variables, such as how frequently people are looking for parking, a few additional variables have been added to the model.

Variable	Definition	Type
----------	------------	------

<i>p</i>	Price of curb-side parking (\$/h)	Union{Real, Distribution}
<i>m</i>	Price of off-street parking (\$/h)	Union{Real, Distribution}
<i>t</i>	Parking duration (h)	Union{Real, Distribution}
<i>f</i>	Fuel cost when cruising (\$/h)	Union{Real, Distribution}
<i>n</i>	People in the car (persons)	Union{Real, Distribution}
<i>v</i>	Value of time (\$/h/person)	Union{Real, Distribution}
<i>ar</i>	Arrival rate (arrivals/min)	Union{Real, Distribution}
<i>cpk</i>	Number of available curb-side spaces	Int64
<i>mint</i>	Minimum parking duration (h)	Real
<i>minc</i>	Minimum time spent coasting (h)	Real
<i>minv</i>	Minimum value of time (\$/h/person)	Real
<i>model_time</i>	Total simulation period (minutes)	Int64
<i>init_occup</i>	Initial curb-side occupancy rate [0,1]	Float64

The arrival rate and other limiting variables need to be specified to ensure convergence. Since the model solves minute-by-minute, an arrival rate of more than 1 arrival per minute is not supported. However, if an arrival rate of more than one vehicle per minute is desired, this can be achieved by increasing the **model_time** and adjusting the other input parameters accordingly to reflect the alteration to the time horizons.

All input parameters into the model are also optional parameters, with each value being assigned a default value and classified in to one of 3 parameter groupings.

Variable	Default Value	Parameter Grouping
<i>p</i>	1.0	Characteristic
<i>m</i>	8.0	Characteristic
<i>t</i>	Distributions.Normal{Float64}(μ=1.0, σ=0.5)	Preference
<i>f</i>	1.0	Characteristic
<i>n</i>	Distributions.Binomial{Float64}(n=3, p=0.8)	Preference
<i>v</i>	Distributions.Normal{Float64}(μ=40.0, σ=5.0)	Preference
<i>ar</i>	Distributions.Bernoulli{Float64}(p=0.2)	Characteristic
<i>cpk</i>	Number of available curb-side spaces	Characteristic
<i>mint</i>	0.1	Characteristic
<i>minc</i>	0.0	Characteristic
<i>minv</i>	0.0	Characteristic
<i>model_time</i>	720	Model
<i>init_occup</i>	0.0	Model

NOTE: The default values represent arbitrary but plausible values.

The model groupings serve no computational purpose, but were included as an aid to organise the parameter inputs in-case of an expansion of the model in future.

Model setup

The parameters are specified by the `init_params()` function, where all input parameters are entered as optional variables. The function returns a tuple containing each of the parameter groupings.

(CharParams, PrefParams, ModelParams)

Each parameter group is another nested tuple containing the relevant variables (see table above). The parameters are then passed to `init_dataframe()` which initialises a dataframe containing the characteristics of each agent, including the arrival time. In cases where a distribution is passed, the characteristics of the agent are independently sampled from the distribution specified. Sample dataframe output below.

5 rows × 9 columns

	t	n	v	p	m	f	ar	cpk	mint
	Float64	Float64	Float64	Float64	Float64	Float64	Float64	Float64	Float64
1	0.618508	3.0	43.1669	1.0	8.0	1.0	0.0	8.0	0.1
2	0.850249	2.0	41.2688	1.0	8.0	1.0	0.0	8.0	0.1
3	1.34617	2.0	38.3695	1.0	8.0	1.0	0.0	8.0	0.1
4	1.09901	1.0	39.5885	1.0	8.0	1.0	0.0	8.0	0.1
5	1.73363	3.0	30.8175	1.0	8.0	1.0	1.0	8.0	0.1

5 rows × 7 columns

	minc	minv	psav	ccost	mct	arrr	tmin
	Float64	Float64	Float64	Float64	Float64	Int64	Float64
1	0.0	0.0	2.39926	172.21	1.0	1	20.5651
2	0.0	0.0	9.40738	119.684	5.0	4	80.6347
3	0.0	0.0	4.33581	79.1385	3.0	5	37.164
4	0.0	0.0	3.71174	70.7016	3.0	11	31.8149
5	0.0	0.0	10.0922	80.431	8.0	15	86.5046

Based on the parameters set, the potential savings from parking on the curb, and the hourly cruising costs are calculated. These values are stored in *psav* and *ccost* respectively and are calculated based on *c* (*c** in the paper). The maximum cruising time and desired cruising duration, in minutes, are given by *mct* and *tmin* respectively. Lastly, *arrr* is the iteration which the agent starts looking for parking which is determined based on the arrival rate parameter, *ar*.

The dataframe and input parameters are then passed to `run_simulation()`, which runs the model over the specified time horizon. The model returns an array with *model_time* rows, populated with a struct of type `ParkState`, where each parkstate contains 8 variables. The struct includes:

Variable	Definition
<code>curb_par_current</code>	Current number of agents parked on the curb
<code>offs_park_current</code>	Current number of agents parked off-street
<code>cruising_curent</code>	Current number of vehicles cruising
<code>curb_park_total</code>	Total number of vehicles parked on the curb
<code>offs_park_total</code>	Total number of vehicles parked off-street
<code>cruising_total_time</code>	Total time spent cruising (hours)
<code>curb_revenue</code>	Revenue for curb-parking provider (\$)
<code>offs_revenue</code>	Revenue for off-street parking provider (\$)

By using the `as_matrix()` function, the simulation output can be converted into a matrix of dimension `[model_time, 8]`. When in matrix format, each column will correspond to one of the variables above, where the order of the columns will correspond with the order of the elements in the struct.

Monte-carlo

When model inputs are passed as distributions, it is more insightful to run a monte-carlo to get a sense of the distribution of the outcome variables of interest. To support this, `mc_simulation()` runs the simulation n_{mc} times. For the monte-carlo simulation, the default output is a 3D array of size `[model_time, 8, n_{mc}]`. Since the dataframe is generated for each model run, the function does not require a dataframe input, only the parameters and the number iterations.

Example

In this example, we'll look at how different pricing policies may impact congestion, and air pollution. For the sake of comparability, we'll be using the figures for the 2020 Honda Civic as a representative city car.

Running a single simulation

First step is to set the characteristics of the desired vehicle, a 2020 Honda Civic in this case.

```

• begin
•   #Setting vehicle emissions, coasting speed, and fuel efficiency (L/hour coasting)
•   co2_kgkm = 0.11
•   nox_kgkm = 1.49e-4
•   coasting_speed_kmh = 8
•   fuel_lh = 0.093*coasting_speed_kmh
•   nothing
• end

```

0.744

```

• fuel_lh

```

Subsequently, the rest of the model parameters can be set. In this case, all the model parameters are called explicitly, however, a parameter does not need to be explicitly set if the default value is desired.

```

• #Setting model parameters

```

```

• pparams, cparams, mparams = init_params(p      = 1.00,
•                                     m          = 13.25,
•                                     t          = Normal(1.5,0.5),
•                                     f          = fuel_lh,
•                                     n          = Binomial(2,0.5),
•                                     v          = Normal(25,5),
•                                     ar         = Bernoulli(0.2),
•                                     cpk       = 8,
•                                     mint       = 0.1,
•                                     minc       = 0.0,
•                                     minv       = 10,
•                                     model_time = 180,
•                                     init_occup = 0.0);
•

```

Once the model parameters have been stored in our parameter variables, we can generate the dataframe and run the simulation.

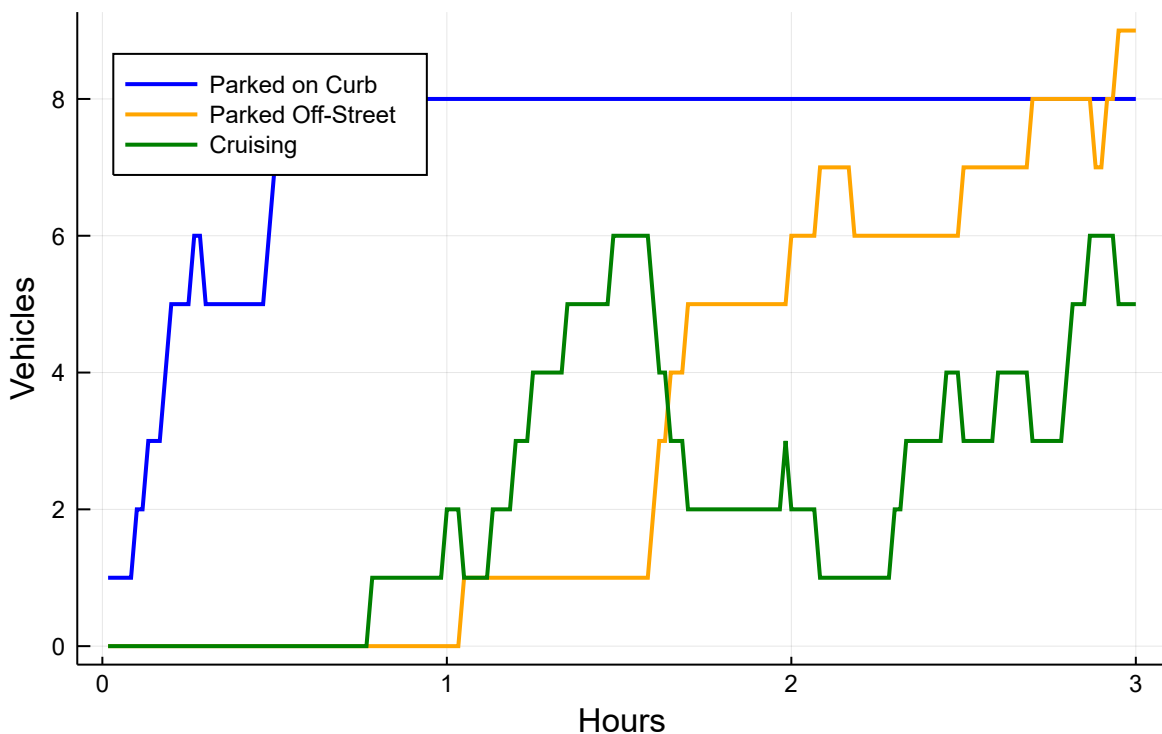
```

• begin
•   #Setting model parameters
•   model_df = init_dataframe(pparams, cparams, mparams)
•
•   #Running simulation
•   model_results = run_simulation(model_df, pparams, cparams, mparams)
•   nothing
• end

```

We can now plot the results with respect to time.

Model Output: Current State



The above graphs shows the state which agents are in along the time-horizon which the model solves for. The green line shows the number of vehicles which are currently cruising to look for parking. Similarly, the blue and orange lines represent the number of vehicles which are currently parked either on the curb or off-street respectively.

Running many simulations

Arguably, running a single simulation is not particularly useful, as we are pulling from several distributions, it could be that we are just getting a tail-case event. Running a monte-carlo, we can get a better sense of what range of results one may expect from the model output.

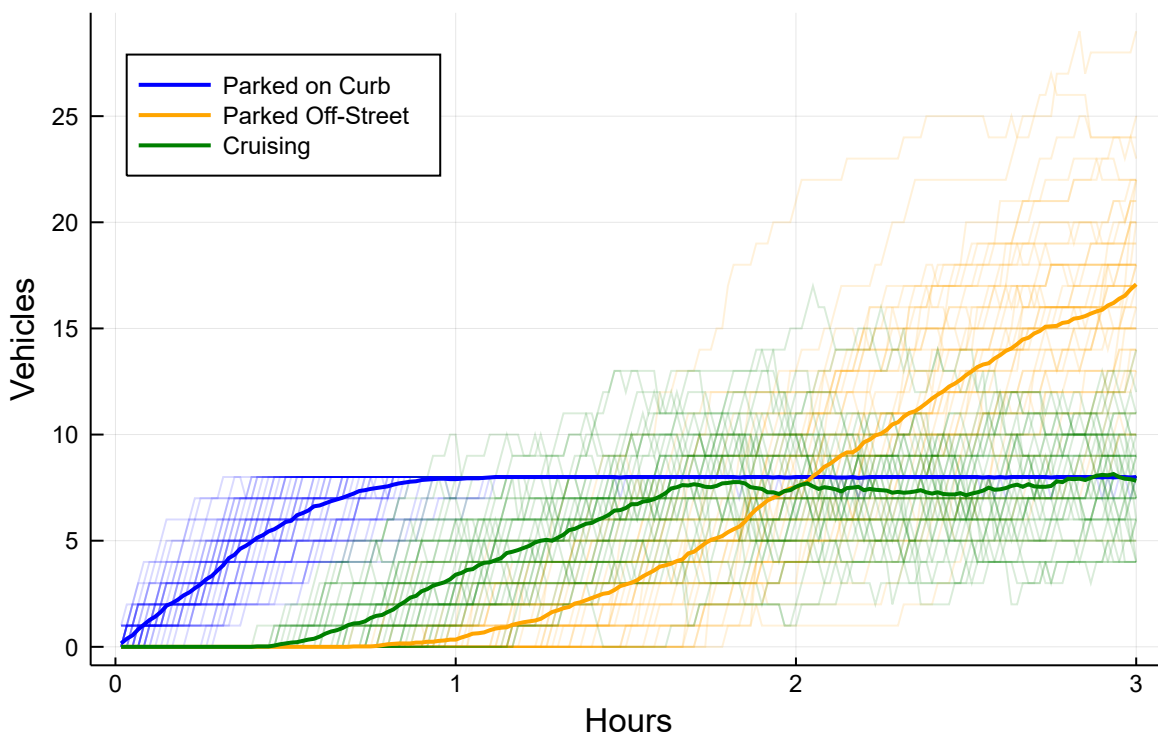
We are going to use the same parameter settings as above, however, we will now also specify the number of times we want to run the simulation.

model iterations: 50

```
• model_results_mc = mc_simulation(pparams, cparams, mparams, n=model_iterations);
```

Plotting the results, we get the following:

Model Output: Current State



curb-side transparency: 0.15

off-street transparency: 0.15

cruising transparency: 0.15

save plot

./cruising_for_parking.png

DPI: 100

Plot title ☒ Model Output: Current State

Autosave ☐

"Plots not currently saving, see REPL for save history"

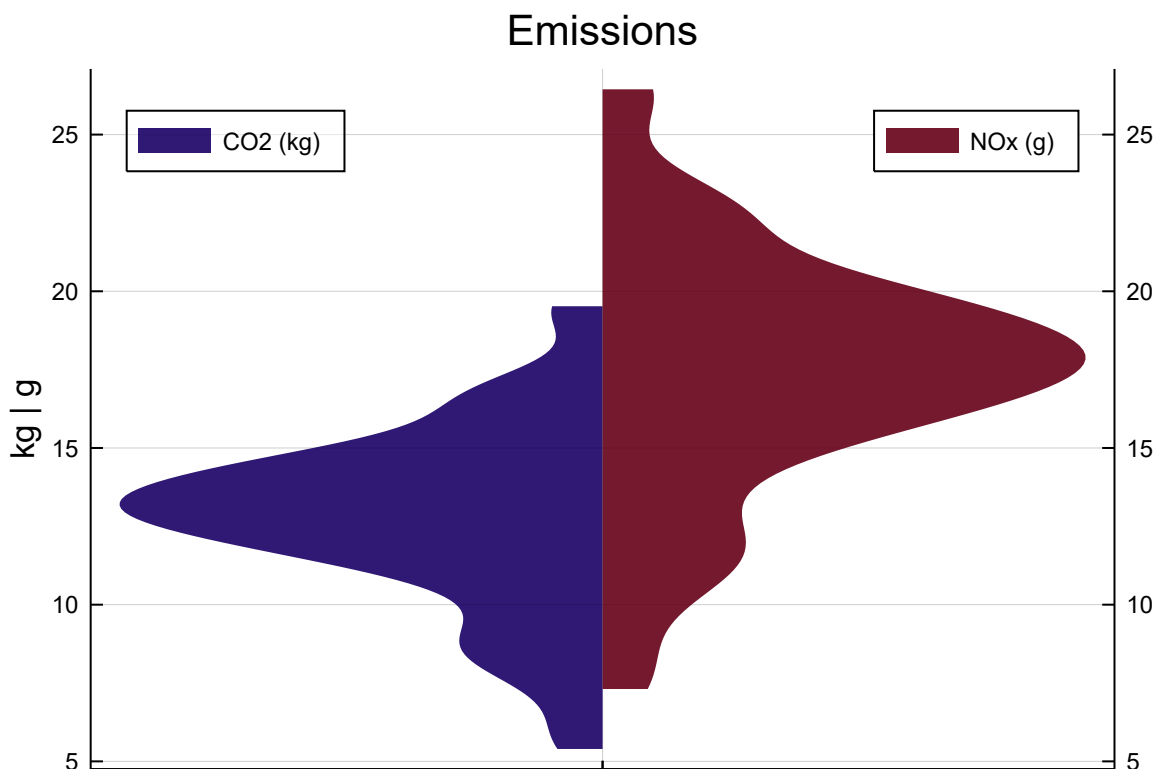
Emissions

Using the vehicle emission data specified earlier (0.11 kg CO₂/km and 0.000149 kg NO_x/km) and assuming an average cruising speed of 8 km/h, we can easily derive the emissions generated from coasting.

```

• begin
•   #Creating the dataframe (function found in appendix)
•   emissions = calculate_emissions(model_results_mc;
•                                   emission_co2=co2_kgkm,
•                                   emission_nox=nox_kgkm) |>
•
•       DataFrame
•
•   #Multiplying nox by 1,000 to be easier to compare to CO2 figures
•   emissions[:nox]*=1000 #g/km
•   emissions[:label] = ""
•   nothing
• end

```



The variance in the emission distribution is determined by variances in cruising time. On the left the CO₂ emissions in kg, and on the right are the NO_x emissions in g.

Distance cruising

On average, 15 hours were spent cruising based on the simulation runs. This equates to the same distance as from Los Angeles to New York 0.03 times, or 116 km. Although some liberties are taken by assuming an infinite space for cruising and that agents not updating their expectation.

Conclusion

Summary

Although at first, the concept of cruising for parking may seem somewhat banal, however, if perverse incentives are set up, it can have large implications for local pollution and congestion. Increasing the price of curb-side parking to align with the price of off-street parking removes the incentive for cruising all together. Increasing the cost of curb-side parking will also raise revenue for the parking provider (typically government) which consequently contribute to the triple-dividend effect.

Limitations

Lastly, I would like to also discuss two main limitations of the current modelling approach. This is by no means meant to be exhaustive, but hopefully brings to light some more questions.

First and foremost, hourly prices seldom work additively. Someone may pay \$10 for parking for 1 hour, and \$12 for parking for 3 hours. Generally, the price curves would be increasing but diminishingly. For the sake of the simulation, this would mean that in the current set-up, people who want to park for a long time are over-estimating the potential savings.

Secondly, agents are myopic and will cruise until the cost of parking on the curb plus the cost of cruising equals the cost of parking off-street. However, in reality, agents likely make decisions based on their expected cruising time, rather than the amount of time realised. If they have already circled the block and all the parking spots are full, they may decide to park off-street without searching any further. At this point, I would also like to point out that the model doesn't account for cyclical road patterns, such as rush hour, which agents may also factor into their expected cruising time.

Appendix

Limitations and improvements

1. Hourly prices seldom work additively over longer parking stays.
2. Agents are currently myopic and do not have a prior expectation of how long it will take to find parking.
3. Street is assumed to be limitless and can host an infinite amount of cruisers
4. Assumes static variables/distributions throughout the modelling period. i.e. doesn't account for rush-hour etc.
5. Assumption regarding that the value of time can be multiplied by the number of people in the vehicle is not necessarily accurate

Function for calculating emissions

```

• function calculate_emissions(results; emission_co2=0.11, emission_nox=1.49e-4,
  coast_speed_kmh=8)
•   #Calculate the total amount of time spend coasting
•   hours_coasting = results[end,6,:]
•
•   #Calculate emissions
•   emissions_co2 = (hours_coasting*coast_speed_kmh)*emission_co2 #Kg of CO2

```

```

• emissions_nox = (hours_coasting*coast_speed_kmh)*emission_nox #Kg of NOx
•
• #Return named tuple with results
• return (co2=emissions_co2, nox=emissions_nox)
• end;

```

Distance calculations

```

• begin
•   dist_la2nyc = 4469
•   average_cruising_time = mean(model_results_mc[end,6,:]) |>
•                           round |>
•                           Int
•   km_travelled = mean(model_results_mc[end,6,:])*coasting_speed_kmh |>
•                  round |>
•                  Int
•   mc_distance_cruised = round(km_travelled/dist_la2nyc, digits=2)
•   nothing
• end

```

Miscellaneous

Be aware that warning messages have been suppressed from showing up in the REPL for this notebook to enable tracking of when plots were saved.

```

• Logging.disable_logging(Logging.Warn);

```

```

• begin
•   if mcplot_title_check
•       mc_plot_title = mcplot_title_field;
•   else
•       mc_plot_title = "";
•   end
•   nothing
• end

```

Packages

```

• begin
•   using StatsPlots, PlutoUI, Distributions, DataFrames, Dates, Logging
•   include("ShoupModel.jl")
•   nothing
• end

```